

Honest Cross-Validation Selects Models That Deploy Better on New Households: Selection Regret in UK Smart-Meter Demand Forecasting

Evgenii Lazarev ^{1,*} 

¹ Independent Researcher; elazarev@gmail.com

Abstract

Cross-validation both estimates error and ranks candidate models, so it decides which model gets deployed. On panel data with repeated entities, shuffled k-fold can rank a model that memorizes entity identity above a simpler model that generalizes, because the same entities appear in training and test folds, and the chosen model then fails on unseen entities. We measure this on next-day household electricity demand from the UK Low Carbon London smart-meter trial. From a fixed eight-candidate grid of k-nearest neighbors, Ridge, and Random Forest, we select once under naive shuffled k-fold and once under group-aware purged cross-validation, then deploy each choice on twelve households held out from selection. Naive cross-validation selects an unconstrained Random Forest that deploys at 2.638 kWh mean absolute error; group-aware cross-validation selects Ridge and deploys at 2.464 kWh, 6.6% lower. Across 30 partitions the group-aware choice deploys better every time (median gap 0.28 kWh), and its model-selection regret against the best-deploying candidate is essentially zero while the naive selector's is not. Plain GroupKFold and a leakage-aware average-load encoding reproduce the effect. The practical rule: match the cross-validation unit to the deployment unit, holding out customers, not rows.

Keywords: cross-validation; model selection; data leakage; load forecasting; smart meters; generalization

1. Introduction

Cross-validation is usually described as a way to estimate how well a model will generalize. In practice it carries a second responsibility that receives less scrutiny: it is the procedure that picks which model gets deployed. An analyst fits a grid of candidates, scores each by cross-validated error, and ships the one with the best score. If the cross-validation score is a faithful proxy for deployment error, the two jobs coincide and the best-scoring model is also the best deploying model. When the proxy is biased in a way that favors one model class over another, the two jobs diverge, and the analyst can ship a model that scored well for the wrong reason.

This divergence has a name. Teodorescu and Obreja Braşoveanu [1] formalize *model-selection regret* as the loss in out-of-sample performance from choosing the model with the best cross-validation score rather than the model with the best out-of-sample score. Studying bankruptcy prediction on independent and identically distributed tabular data, they find that k-fold cross-validation is a valid selection rule on average, but that roughly 67% of the variance in regret is attributable to the particular train and test split rather than to the selection procedure, which they read as an irreducible uncertainty that more

Received:

Revised:

Accepted:

Published:

Copyright: © 2026 by the author.

Submitted to *Forecasting* for possible open access publication under the terms and conditions of the [Creative Commons Attribution \(CC BY\)](#) license.

elaborate cross-validation cannot remove. They also report a more pointed result and leave it open: decisions *between model classes* taken on cross-validation performance are, more often than not, incorrect, and this warrants further investigation. The present paper takes up exactly that open question, in a setting where the data are not independent and identically distributed but grouped and time-indexed, and where the regret source is not irreducible but a removable form of leakage.

Leakage is a well-documented cause of inflated performance estimates in machine learning [2], and it is especially costly when a model is selected and reported after many iterations, so that the optimistic score becomes part of the evidence rather than a private development error [3]. Kapoor and Narayanan [4] catalogue leakage as a recurring driver of a reproducibility crisis across machine-learning-based science, and Roberts et al. [5] document how leakage and improper validation invalidated a large body of published models. A common and under-recognised form is entity leakage in panel data, where rows from the same group appear in both training and test folds. In financial machine learning, López de Prado [6] introduced purging and embargoing to guard against leakage from overlapping labels and serial dependence, and group-disjoint folds to evaluate generalization to unseen entities. Recent work continues to document how validation design alone can swing reported accuracy in time-series forecasting [7]. Most of this literature measures leakage as *bias in the score*. We measure something operationally sharper: whether leakage changes *which model is selected*, and what that wrong selection costs at deployment.

Smart-meter demand forecasting is a natural place to study this. The data are a panel: many households, each observed over time. A common deployment target is a new customer with little or no history, for whom the relevant generalization unit is the household, not the calendar row. Yet the convenient default, shuffled k-fold, scatters a household's rows across training and test folds, so a model can learn a per-household baseline during selection and look excellent, then fail on customers it never saw.

This paper makes four contributions. First, we give a concrete demonstration on real energy data that the cross-validation scheme changes the deployed model class, not merely the reported error. Second, we quantify model-selection regret against the best-deploying candidate on truly held-out households and show that the group-aware selector is at or near zero regret across 30 partitions while the naive selector is not. Third, we isolate the mechanism from two sides: an ablation that removes the entity feature and erases the gap, and an innocent-feature variant in which a leakage-aware encoding of the customer's average load, the kind of feature a careful practitioner would use, reproduces the gap. Fourth, the entire study is reproducible from open code over a public dataset. We use the scikit-learn-compatible `purgedcv` package [8,9] for the group-aware split, and we confirm that the ordinary scikit-learn `GroupKFold` produces the same selection, so the result rests on group-aware validation as a principle, not on a particular implementation.

2. Materials and Methods

2.1. Data

We use the Low Carbon London smart-meter dataset released by UK Power Networks through the London Datastore [10]. The analysis uses a cached 60-household subset of half-hourly consumption, the same subset shipped with the package examples, drawn from the Standard-tariff households with a full year of 2013 readings. The series is restricted to the 2013 calendar year and aggregated to daily consumption per household in kWh. The prediction target is a household's daily consumption; the task is partially predictable, because daily load has strong day-of-week and short-lag structure but also large idiosyncratic variation. This subset is fixed throughout; we do not claim it is a random sample of the

full London population, and we return to that point with a full-population benchmark in Section 3.6.

Households are the unit of deployment. For each random seed we permute the 60 households and assign 48 to a selection set and the remaining 12 to a deployment set that is never seen during model selection. With seed 0 the selection set contains 16,090 daily observations from 48 households and the deployment set contains 4,122 observations from 12 held-out households.

2.2. Features

Each row carries eleven features: the previous day's consumption (`lag_1`), the same weekday one week earlier (`lag_7`), a one-hot encoding of the day of week (`dow_0` through `dow_6`), and the household identifier mapped to an integer (`hh_int`). The identifier is an arbitrary ordinal: households are numbered by sorting their string identifiers, so the integer value carries no meaningful order. It is included on purpose. Inside shuffled k-fold, where the same household appears in both training and test folds, a sufficiently flexible model can split on `hh_int` and memorize each household's baseline level, which lowers the cross-validation error without improving generalization. On a genuinely unseen household the identifier takes a value the model never trained on and carries no information.

We use the raw identifier because it makes the mechanism unambiguous, but the point is more general than this one feature. Any feature correlated with entity identity opens the same route: a categorical customer or meter identifier, a tariff or postcode code, or an engineered per-customer statistic. Such features are routine in applied load-forecasting pipelines, and the failure they enable is invisible to shuffled cross-validation, which scores well precisely because the leaked information is present in both folds. We treat the single identifier as a controlled probe of that broader class and acknowledge its artificial simplicity in the limitations.

2.3. Candidate grid

The candidate grid spans eight models from overfitting-prone to constrained: k-nearest neighbors with $k \in \{1, 5, 50\}$, Ridge regression with $\alpha \in \{0.01, 1, 100\}$, and Random Forest regressors with `max_depth` of `None` (unconstrained, free to split deeply on `hh_int`) and 4 (shallow). The k-nearest-neighbors and Ridge candidates are standardized through a scikit-learn pipeline. The grid is fixed in advance and identical across both cross-validation schemes, so any difference in the selected model is caused by the validation design and nothing else.

2.4. Cross-validation schemes

Two five-fold schemes are compared. The naive scheme is `KFold(n_splits=5, shuffle=True)`, which ignores household identity and scatters rows at random, so training and test folds routinely share households. The honest scheme is `PurgedGroupKFold` from `purgedcv`, which keeps whole households together, so every test fold contains households absent from the corresponding training fold. This mirrors the deployment condition exactly. Because the labels here are instantaneous daily values rather than forward-horizon intervals, we set the evaluation time equal to the prediction time, so the group-aware splitter performs a pure household-disjoint split with no time-purging. This isolates the group-leakage effect from any overlapping-label effect.

In this instantaneous-label setting the purged group split reduces to an ordinary group-disjoint split, which scikit-learn's `GroupKFold` also provides. We therefore run a third scheme, `GroupKFold(n_splits=5)` grouped by household, as a control. Its purpose is honesty: if the simple group splitter recovers the same selection, then the result is attributable to group-aware validation as a principle and not to any feature unique to

purgedcv. The added value of purgedcv over GroupKFold appears only when labels also span future horizons, so that purging and grouping must be combined; that regime is exercised by the full-population benchmark and the other examples in the repository rather than by this experiment.

2.5. Selection-regret protocol

Following Teodorescu and Obreja Braşoveanu [1], define the model-selection regret of a selector that picks model M^{CV} as

$$\delta = \text{OOS}(M^{\text{CV}}) - \text{OOS}(M^{\text{OOS}}), \quad (1)$$

where $\text{OOS}(\cdot)$ is deployment mean absolute error on the 12 held-out households and M^{OOS} is the candidate in the grid with the lowest such error. Regret is zero when the selector picks the best-deploying candidate and positive otherwise. To compute it we record the deployment error of *every* candidate, identify M^{OOS} , and evaluate Equation (1) separately for the naive and the honest selector. We also report the directly practitioner-visible quantity, the naive deployment error minus the honest deployment error, which is positive when the honest selector deploys better; this difference equals the naive selector's regret whenever the honest selector picks the best-deploying candidate.

The protocol is identical for every scheme. Each candidate is scored by five-fold cross-validated mean absolute error on the selection set. The lowest-error candidate is selected, retrained on the full selection set, and evaluated on the deployment households by mean absolute error and coefficient of determination R^2 . To move beyond a single realization, the whole pipeline, including the household partition, the shuffle, and the Random Forest seeds, is repeated for 30 seeds.

A statistical caveat governs how we read those 30 repetitions. Each seed is a different 48/12 partition of the *same* fixed 60-household sample. These are correlated resamples without replacement, not independent draws from the household population, so we summarise them with descriptive statistics (median, interquartile range, range) and a win count rather than with a population confidence interval, which would understate the true uncertainty. Population-level inference is left to the independent 20-subsample benchmark of Section 3.6.

2.6. Ablation

To test that the household identifier is the mechanism, the experiment is repeated with `hh_int` removed from the feature set. If the identifier was the only route to leakage, both selectors should now pick the same model and the deployment gap should vanish.

2.7. Innocent-feature variant

The raw identifier is a deliberately blunt probe. To check that the effect is not an artifact of feeding a raw identifier to a tree, we run a second variant that replaces `hh_int` with the feature a careful practitioner would actually engineer: the customer's average daily consumption, the standard baseline-load predictor. It is built without naive target leakage. The per-household mean of the target is computed by a leakage-aware encoder fit only on the training side of each fold, and a household absent from the training side, the cold-start condition of a genuinely new customer, is imputed with the global training mean. Each candidate becomes a pipeline of this encoder followed by the estimator, so the feature is recomputed correctly inside every fold. Under shuffled folds a household still appears on both sides, so its test rows receive its own mean; under group-aware folds the test households are imputed with the global mean, exactly as at deployment. This

variant uses the same grid, schemes, deployment protocol, and 30 partitions as the main experiment, and is reproduced by `examples/selection_regret_lcl_targetenc.py`.

2.8. Computational environment

The selection-regret experiments reported here, including the 30-partition sweep and the innocent-feature variant, were produced with `purgedcv` 0.0.11, Python 3.12.7, `numpy` 1.26.4, `pandas` 2.3.3, `scikit-learn` 1.8.0, and `scipy` 1.17.1 on macOS 26.3.1 (arm64). They are reproduced by the script `examples/selection_regret_lcl_seeds.py`, the innocent-feature variant `examples/selection_regret_lcl_targetenc.py`, and the companion notebook `examples/selection_regret_lcl.ipynb`, which reproduces the single-seed walk through and the headline figure. The full-population benchmark of Section 3.6 was run separately, on a larger compute job and an earlier software environment, so its figures should be read as independent corroboration rather than as output of the environment above.

3. Results

3.1. The two schemes select different model classes

Table 1 shows the five-fold cross-validated mean absolute error of every candidate under both schemes, for seed 0. Under naive shuffled k-fold the unconstrained Random Forest posts the lowest error, 1.6520 kWh, because it can read each household’s baseline off the identifier feature and the same households appear in the test fold. Under group-aware cross-validation that advantage disappears: the unconstrained Random Forest rises to 2.0018 kWh, and the lowest-error candidate is Ridge with $\alpha = 0.01$ at 1.7094 kWh. The two schemes therefore select different model classes from the same grid on the same data.

Table 1. Five-fold cross-validated mean absolute error (kWh) per candidate on the selection set, seed 0. The selected model in each column is the row with the lowest error.

Candidate	Naive shuffled k-fold	Group-aware (honest)
kNN ($k = 1$)	2.1175	2.7864
kNN ($k = 5$)	1.7413	2.2254
kNN ($k = 50$)	1.7603	2.2544
Ridge ($\alpha = 0.01$)	1.6913	1.7094
Ridge ($\alpha = 1$)	1.6914	1.7095
Ridge ($\alpha = 100$)	1.6948	1.7136
RF (max_depth=None)	1.6520	2.0018
RF (max_depth=4)	1.7197	1.8649
Selected	RF (max_depth=None)	Ridge ($\alpha = 0.01$)

3.2. The honest selection deploys better on unseen households

Table 2 reports what happens when every candidate is retrained on the full selection set and evaluated on the 12 held-out households, for seed 0. The best-deploying candidate M^{OOS} is Ridge with $\alpha = 100$ at 2.4573 kWh; the three Ridge variants deploy within 0.007 kWh of one another, so the regularized linear family is effectively the deployment optimum. The Random Forest chosen by naive cross-validation deploys at 2.6376 kWh, a regret of 0.1803 kWh by Equation (1). The Ridge model chosen by group-aware cross-validation deploys at 2.4641 kWh, a regret of only 0.0068 kWh: the honest selector lands essentially on the deployment optimum. The practitioner-visible gap between the two selections is 0.1735 kWh, a 6.6% reduction, with the higher coefficient of determination (0.7834 versus 0.7587). Figure 1 shows the contrast.

Table 2. Deployment mean absolute error (kWh) on the 12 held-out households for every candidate, seed 0. M^{OOS} marks the best-deploying candidate; regret is each selector's deployment error minus that of M^{OOS} .

Candidate	Deployment MAE (kWh)	Note
kNN ($k = 1$)	3.6393	
kNN ($k = 5$)	2.7378	
kNN ($k = 50$)	2.5440	
Ridge ($\alpha = 0.01$)	2.4641	honest pick (regret 0.0068)
Ridge ($\alpha = 1$)	2.4640	
Ridge ($\alpha = 100$)	2.4573	M^{OOS} (best)
RF (max_depth=None)	2.6376	naive pick (regret 0.1803)
RF (max_depth=4)	2.5769	
Naive minus honest deployment MAE		+0.1735 (+6.6%)




Figure 1. Deployment mean absolute error on the 12 held-out households for the model each scheme selects, seed 0. Same grid, same data, same held-out households; the honest selection deploys at 6.6% lower error.

3.3. The advantage holds across 30 random partitions

A single household partition is one draw, and Teodorescu and Obreja Braşoveanu [1] warn that the partition itself is the dominant source of regret variance. We therefore repeat the full pipeline over 30 partitions, reading them as correlated resamples of the fixed 60-household sample as explained in the methods. The selection does not waver. Naive cross-validation selects the unconstrained Random Forest in all 30 partitions, and group-aware cross-validation selects Ridge with $\alpha = 0.01$ in all 30. The ordinary GroupKFold control selects the same model as the purged splitter in all 30, which puts the effect on group-aware validation itself rather than on a specific package. In 24 of the 30 partitions the honest pick is also the best-deploying candidate M^{OOS} ; in the other six a neighbouring Ridge setting beats it by at most 0.007 kWh.

The two selectors are therefore very different in regret. The honest selector's regret has a median of 0.0000 kWh (interquartile range 0.0000 to 0.0000): it is at the deployment optimum in almost every partition. The naive selector's regret has a median of 0.2831 kWh (interquartile range 0.1831 to 0.5157). The practitioner-visible gap, naive minus honest deployment error, is positive in all 30 partitions, with a median of 0.2831 kWh (interquartile range 0.1826 to 0.5157, range 0.1084 to 1.7851) and a median relative reduction of 13.3% (interquartile range 8.1% to 22.5%). Table 3 summarizes, and Figure 2 shows the paired deployment errors and the distribution of the gap, which lies entirely on the positive side of zero.

Table 3. Robustness over 30 correlated partitions of the 60-household sample. Spreads are descriptive (interquartile range, IQR), not population confidence intervals; see the statistical caveat in the methods. Regret is computed against the best-deploying candidate by Equation (1).

Quantity	Value
Naive selects RF (max_depth=None)	30 / 30
Honest selects Ridge ($\alpha = 0.01$)	30 / 30
GroupKFold selects same as honest	30 / 30
Honest pick is best-deploying M^{OOS}	24 / 30
Honest deploys better than naive	30 / 30
Naive selector regret, median (IQR), kWh	0.2831 (0.1831 to 0.5157)
Honest selector regret, median (IQR), kWh	0.0000 (0.0000 to 0.0000)
Naive minus honest MAE, median (IQR), kWh	0.2831 (0.1826 to 0.5157)
range, kWh	0.1084 to 1.7851
Naive minus honest MAE, median relative	13.3%




Figure 2. Left: deployment mean absolute error for the naive and honest selections in each of the 30 partitions; every paired line falls. Right: distribution of the naive minus honest deployment-error gap; all of the mass is positive, meaning the honest selection deploys better in every partition.

3.4. Removing the identifier erases the gap

When the household identifier `hh_int` is removed from the feature set, both schemes select Ridge with $\alpha = 0.01$, and both deploy at the same mean absolute error of 2.4676 kWh on the held-out households. The regret difference falls to zero. This pins the mechanism precisely: the gap was not a property of Random Forests in general, nor of the cross-validation arithmetic, but of one entity-identifier feature that shuffled cross-validation allowed a flexible model to memorize. Group-aware cross-validation refuses that information during selection, and the selection then matches what deployment will look like.

3.5. An innocent, best-practice feature reproduces the effect

A reviewer may object that a skilled practitioner would never feed a raw identifier to a tree, so the gap could be an artifact of a contrived feature. The innocent-feature variant answers this. Replacing the identifier with a leakage-aware target-mean encoding of the household, the customer's average daily consumption built per fold with cold-start imputation, reproduces the result. The gap is larger, not smaller. Across the 30 partitions, naive cross-validation again selects the unconstrained Random Forest in all 30, while group-aware cross-validation selects Ridge in 29 and the shallow Random Forest once. The honest selector picks the best-deploying candidate in all 30 partitions, so its median regret is 0.0000 kWh; the naive selector's regret has a median of 0.3645 kWh (interquartile range 0.2998 to 0.4285). The honest selection deploys better in all 30 partitions, with a median gap of 0.3645

kWh (range 0.1754 to 0.8028) and a median relative reduction of 15.7%. On the seed-0 partition the naive Random Forest deploys at 3.1048 kWh and the honest Ridge at 2.5011 kWh, a gap of 0.6038 kWh. The mechanism is therefore not specific to a raw identifier: any feature that encodes entity identity, even one engineered with correct per-fold hygiene, leaks under shuffled folds because the entity sits on both sides of the split, and the leak is invisible to the cross-validation score that selects the model.

3.6. Household structure carries a large generalization penalty at scale

The single-grid study above uses one fixed 60-household sample. A full-population benchmark on the same corpus confirms that the household structure carries a substantial generalization penalty well beyond that sample, and it does so with genuinely independent subsamples, which the 30 partitions deliberately are not. Scanning 167,932,474 raw rows and drawing 20 independent seeded subsamples of 60 households each from the 4,284 eligible households, scoring on entirely unseen households is 6.03% worse in weighted absolute percentage error than a pooled temporal estimate (95% confidence interval 4.93% to 7.12%), whereas the purely temporal leakage gap is only 1.60%. The metric there is weighted absolute percentage error rather than the mean absolute error used above, and the estimator is a single gradient-boosted model rather than a selection grid, so the two studies are complementary rather than directly comparable; what they share is the conclusion that the dominant operational risk in this corpus is household generalization rather than temporal leakage, which is exactly the risk that group-aware selection is built to measure.

4. Discussion

The result has a simple cause. Shuffled cross-validation lets a model see, during scoring, information it will not have at deployment time, namely the identity and baseline level of households that will also appear in the test fold. A deep Random Forest exploits this by learning per-household leaf values, which makes its cross-validation error look the best in the grid. On households it has never seen, the identifier is meaningless and the memorized structure is dead weight, so the deep forest falls behind a regularized linear model. Group-aware cross-validation removes the leaked information during selection, so its ranking reflects the deployment condition and it selects the model that actually generalizes.

Teodorescu and Obreja Braşoveanu [1] conclude that on independent and identically distributed tabular data the train and test split dominates regret variance, so more elaborate cross-validation buys little on average. Our finding is not a contradiction but a complement that depends on the data regime. On their i.i.d. data the regret source is largely irreducible. On grouped, time-indexed data the regret source is reducible: it is entity leakage, and a validation scheme matched to the deployment unit removes it. We also answer their explicit open question. They observe that decisions between model *classes* taken on cross-validation performance are often wrong and call for further investigation; here the cross-class decision (Ridge versus Random Forest) is wrong under naive cross-validation in every one of 30 partitions, while under group-aware cross-validation the chosen model deploys better than the naive choice in all 30 and is at or near the deployment optimum.

This is the leakage that Kapoor and Narayanan [4] and Roberts et al. [5] identify as a recurrent cause of irreproducible models, here in its entity-leakage form on panel data and traced to its effect on *selection* rather than on the reported score alone. That load models generalize poorly to unseen households is itself established [11]. What we add is that the standard validation default actively rewards the model that generalizes worst, and that a group-aware split repairs the selection. We expect the same pattern wherever a repeated-entity panel is validated by shuffled folds, for example patients, devices, stores,

or sensors, although we have demonstrated it on one dataset and treat that broader claim as a hypothesis rather than an established result.

The actionable rule is to match the cross-validation unit to the deployment unit. If a demand forecaster will serve customers who were present during training, a pooled or temporal split may be appropriate. If it will serve new customers, selection must hold out customers, because that is the question deployment will ask. The choice is not a tuning detail; it determines which model class is shipped. The `purgedcv` package makes the group-aware split a drop-in scikit-learn cross-validator, so adopting the honest selection costs one line of code.

The study uses one dataset, one forecasting target at daily resolution, and a deliberately compact grid; the magnitudes will differ on other corpora, horizons, and richer grids, and a fuller study would sweep these axes. The household identifier is an intentional and deliberately simple trap. A skilled practitioner would not feed a raw identifier to a tree, and to that extent the specific feature is artificial; the defensible claim is the general one, that any entity-correlated feature opens the same route and that shuffled cross-validation gives no warning when it does. The ablation confirms this by erasing the gap once the feature is removed, and the innocent-feature variant confirms it from the other side: a leakage-aware target encoding of the customer’s average load, a feature a careful practitioner would use, reproduces the gap rather than removing it. The robustness sweep covers 30 partitions of a single 60-household sample; as noted, these are correlated resamples, so the descriptive spreads are not population confidence intervals, and the independent-subsample benchmark uses a different metric and estimator. Finally, group-aware validation addresses entity leakage; it does not by itself remove other leakage routes such as globally fit preprocessing or future-looking feature engineering.

5. Conclusions

On a real, partially predictable demand-forecasting task, the cross-validation scheme used for model selection decided which model class was deployed. Naive shuffled k-fold consistently selected an unconstrained Random Forest that memorized household identity and deployed worse on new households, while group-aware purged cross-validation consistently selected Ridge and deployed better, by 6.6% on the headline split and in every one of 30 random partitions. The effect is removable, not irreducible: it is entity leakage, and a validation unit matched to the deployment unit removes it. For demand forecasting aimed at new customers, and for panel forecasting more broadly, validation should be performed at the granularity of the entity the model will actually face.

Supplementary Materials: The following supporting information is provided with this article: the per-partition selection-regret results `selection_regret_lcl_seeds.csv` (main experiment) and `selection_regret_lcl_targetenc.csv` (innocent-feature variant), and the full-population benchmark summary `lcl_full_benchmark_summary.md`.

Author Contributions: Conceptualization, methodology, software, validation, formal analysis, investigation, data curation, writing—original draft, writing—review and editing, and visualization were all carried out by E.L. The author has read and agreed to the published version of the manuscript.

Funding: This research received no external funding.

Institutional Review Board Statement: Not applicable.

Informed Consent Statement: Not applicable.

Data Availability Statement: The Low Carbon London smart-meter data are publicly available from UK Power Networks and the London Datastore [10]. The cross-validation splitter is provided by the open-source `purgedcv` package, available at <https://github.com/eslazarev/purged-cross-validation>

and on PyPI as `purgedcv` [8]; the software is MIT licensed and archived on Zenodo with DOI [10.5281/zenodo.20312695](https://doi.org/10.5281/zenodo.20312695). The repository contains the experiment scripts and notebook that regenerate every reported number and figure from the public data: `examples/selection_regret_lcl_seeds.py`, the innocent-feature variant `examples/selection_regret_lcl_targetenc.py`, and the single-seed walk-through `examples/selection_regret_lcl.ipynb`. The cached 60-household extract (about 68 MB) and the derived seed-level result files are not redistributed in the repository owing to size; they are regenerated by these scripts from the public dataset, and the small derived result tables are provided as Supplementary Materials.

Acknowledgments: During the preparation of this manuscript the author used OpenAI Codex and ChatGPT (GPT-5 family) and Anthropic Claude (Claude 4 family) for code review, refactoring suggestions, and copy-editing. The author has reviewed and edited the output and takes full responsibility for the content of this publication; no numerical claim was accepted without executable verification against the experiment script.

Conflicts of Interest: The author declares no conflicts of interest.

References

- Teodorescu, V.; Obreja Braşoveanu, L. Assessing the Validity of k-Fold Cross-Validation for Model Selection: Evidence from Bankruptcy Prediction Using Random Forest and XGBoost. *Computation* **2025**, *13*, 127. <https://doi.org/10.3390/computation13050127>.
- Kaufman, S.; Rosset, S.; Perlich, C.; Stitelman, O. Leakage in Data Mining: Formulation, Detection, and Avoidance. *ACM Transactions on Knowledge Discovery from Data* **2012**, *6*, 1–21. <https://doi.org/10.1145/2382577.2382579>.
- McDermott, M.B.A.; Wang, S.; Marinsek, N.; Ranganath, R.; Foschini, L.; Ghassemi, M. Reproducibility in Machine Learning for Health Research: Still a Ways to Go. *Science Translational Medicine* **2021**, *13*. <https://doi.org/10.1126/scitranslmed.abb1655>.
- Kapoor, S.; Narayanan, A. Leakage and the Reproducibility Crisis in Machine-Learning-Based Science. *Patterns* **2023**, *4*, 100804. <https://doi.org/10.1016/j.patter.2023.100804>.
- Roberts, M.; Driggs, D.; Thorpe, M.; et al. Common Pitfalls and Recommendations for Using Machine Learning to Detect and Prognosticate for COVID-19 Using Chest Radiographs and CT Scans. *Nature Machine Intelligence* **2021**, *3*, 199–217. <https://doi.org/10.1038/s42256-021-00307-0>.
- López de Prado, M. *Advances in Financial Machine Learning*; Wiley: Hoboken, NJ, 2018. Purging and embargoing: chapter 7. Combinatorial Purged Cross-Validation: chapter 12.
- Albelali, S.; Ahmed, M. Hidden Leaks in Time Series Forecasting: How Data Leakage Affects LSTM Evaluation Across Configurations and Validation Strategies, 2025, [arXiv:cs.LG/2512.06932]. <https://doi.org/10.48550/arXiv.2512.06932>.
- Lazarev, E. `purgedcv`: scikit-learn-compatible purged and combinatorial cross-validation for time-series and panel machine learning. Python package, <https://github.com/eslazarev/purged-cross-validation>, 2026. Software concept DOI; MIT license., <https://doi.org/10.5281/zenodo.20312695>.
- Pedregosa, F.; Varoquaux, G.; Gramfort, A.; Michel, V.; Thirion, B.; Grisel, O.; Blondel, M.; Prettenhofer, P.; Weiss, R.; Dubourg, V.; et al. Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research* **2011**, *12*, 2825–2830.
- UK Power Networks. SmartMeter Energy Consumption Data in London Households (Low Carbon London). London Datastore, 2014. Dataset `smartmeter-energy-use-data-in-london-households`; open terms.
- Pirbazari, A.M.; Farmanbar, M.; Chakravorty, A.; Rong, C. Short-Term Load Forecasting Using Smart Meter Data: A Generalization Analysis. *Processes* **2020**, *8*, 484. <https://doi.org/10.3390/pr8040484>.